Reviewer Memo: December 15, 2025 Update

From Conditional to Unconditional Hodge Closure

Jonathan Washburn
Recognition Science, Recognition Physics Institute
jon@recognitionphysics.org

December 15, 2025

Executive Summary

The manuscript hodge-SAVE-dec-12-handoff.tex has been updated from a **conditional** proof (pending the "microstructure/gluing" step) to an **unconditional** proof of the Hodge Conjecture for rational (p, p) classes on smooth projective Kähler manifolds.

The previous version you reviewed (hodge-fix-dec-8-old.tex) established the calibration–coercivity framework and the quantitative approximation to the calibrated cone. That machinery remains unchanged. What has been added is the *realization/microstructure* step that was previously flagged as open.

Key change: The proof now includes a complete "corner-exit vertex-template" construction that manufactures ψ-calibrated holomorphic complete intersections with controlled geometry, enabling the flat-norm gluing estimate $\mathcal{F}(\partial T^{\mathrm{raw}}) = o(m)$.

1 What Was Missing Before (and Why)

The previous manuscript reduced the Hodge Conjecture to the following:

> **Microstructure/Gluing Theorem (informal).** For every smooth closed cone-valued (p, p)-form β representing an effective rational Hodge class, and for m large and mesh h small, there exists a construction of ψ-calibrated holomorphic complete-intersection pieces whose sum T^{raw} satisfies:
>
> 1. T^{raw} approximates $m\beta$ cellwise in the tangential/Young-measure sense, and
> 2. $\mathcal{F}(\partial T^{\mathrm{raw}}) = o(m)$.

The obstruction was that while we could build *local* holomorphic pieces inside each cell, we could not control their *boundary mismatch* across cell interfaces well enough to guarantee the flat-norm bound (ii).

The difficulty is combinatorial/geometric: on each interior face $F = Q \cap Q'$, the pieces from Q and Q' contribute boundary slices that must nearly cancel. Without careful coordination, these mismatches can accumulate to $O(m)$ rather than $o(m)$.

2 The Solution: Corner-Exit Vertex Templates

The key insight is to use a **vertex-anchored, corner-exit** construction:

1. **Corner localization.** Each holomorphic sliver is anchored at a vertex v of a cube Q and has footprint contained in $B(v, c_0 h)$ with $c_0 < 1$. This automatically prevents the sliver from reaching any face *not* incident to v.
2. **Designated exit faces.** The footprint is a uniformly fat k-simplex ($k = 2n - 2p$) whose $k+1$ facets lie on exactly $k+1$ coordinate faces incident to v (the "exit faces"). This gives a clean "if and only if" characterization of which faces a sliver meets.
3. **Prefix-based activation.** At each vertex, slivers are organized into an ordered template $(P_{v,1}, P_{v,2}, \ldots)$. To realize N slivers at vertex v, we activate the prefix $\{P_{v,a}\}_{a \le N}$. Because the template is the same on both sides of a shared vertex, slow variation of counts ($|N_{Q,v} - N_{Q',v}| \lesssim h \cdot N$) means mismatches are confined to the "tail" of the template.
4. **No heavy tails.** Because all slivers in a template have comparable mass (by equal footprint geometry), the tail mismatch is an $O(h)$ fraction of the total—not a rare heavy piece that dominates.

This reduces the "global face consistency" problem to a *local geometric* problem: can we actually manufacture holomorphic slivers with the required corner-exit footprint?

3 New Lemmas and Propositions (Manuscript References)

The following are the main new results added to achieve unconditional closure:

A. Corner-Exit Geometry (Euclidean Model)

- `lem:ball-excludes-faces` (ball locality excludes nonincident faces): If $E \subset B(v, c_0 h)$ with $c_0 < 1$, then E cannot meet any face not incident to v.
- `lem:corner-simplex-hits-designated-faces` (fat corner simplex hits designated faces): A uniformly fat k-simplex with the corner-exit structure meets each designated exit face in a genuine $(k-1)$-patch.
- `lem:corner-simplex-face-mass` (uniform per-face boundary mass): Each facet has $(k-1)$-mass comparable to (sliver mass)$^{(k-1)/k}$.
- `lem:small-graph-distortion` (graph area distortion): A C^1 graph with slope $\le \varepsilon$ distorts k- and $(k-1)$-areas by $1 + O(\varepsilon^2)$.

B. Complex Corner-Exit Templates

- `lem:complex-corner-exit-template` (explicit complex example): Constructs a concrete complex $(n-p)$-plane family with the corner-exit property, showing existence.
- `lem:corner-exit-template-open` (quantitative template family): For a complex plane satisfying quantitative nondegeneracy bounds on its coefficients, there exists a corner-exit translation template with uniform fatness. The key point: one can choose vertex and exit-face set, and there is a $(2p-1)$-parameter box of translations giving *identical* footprints.

- `prop:corner-exit-template-net` (robust templates for a finite net): For any ε_h-net of calibrated directions, one can perturb to ensure *every* direction admits corner-exit templates, with uniform constants over the finite net.

C. Holomorphic Realization

- `lem:global-graph-contraction` (contraction criterion for global graphs): If a holomorphic map $F(u, w)$ satisfies $|\partial_w F - I| \le \eta$ uniformly on a product domain, then $\{F = 0\}$ is a single-sheet C^1 graph $w = g(u)$ on the entire domain.

- `lem:bergman-affine-approx-hormander` (Bergman-scale affine approximation): Using cutoff + Hörmander $\bar{\partial}$-solving, one can construct global holomorphic sections whose coefficient functions are uniformly C^1-close to any prescribed affine-linear model on a ball of radius $R/\sqrt{m}$. The error is $O(e^{-cm})$—exponentially small.

- `prop:cell-scale-linear-model-graph` (cell-scale single-sheet graphs): Combining the above, holomorphic complete intersections at Bergman scale ($h \lesssim m^{-1/2}$) are single-sheet graphs on entire cells.

- `prop:holomorphic-corner-exit-g1g2` (holomorphic slivers inherit corner-exit geometry): Holomorphic small-slope graphs over fat corner-exit footprints satisfy (G1-iff) and (G2).

D. Global Assembly

- `prop:vertex-template-mass-matching` (L2: mass budget matching): Nearest-integer rounding of prefix lengths matches cell mass budgets with relative error $O(1/N) + O(\varepsilon^2)$.

- `prop:vertex-template-face-edits` (face-level $O(h)$ edits): Slow variation of counts implies unmatched boundary mass on each face is an $O(h)$ fraction.

- `prop:global-coherence-all-labels` (B1: global coherence across all direction labels): Packages the full execution: choose corner-exit-admissible direction nets, Lipschitz weights, slow-variation integer counts, and cohomology discrepancy rounding; sum labelwise flat-norm bounds to get $\mathcal{F}(\partial T^{\text{raw}}) = o(m)$.

4 Proof Strategy Overview

The complete proof now has the following structure:

1. **Signed decomposition (unchanged).** Any rational Hodge class γ decomposes as $\gamma = \gamma^+ - \gamma^-$ with both effective. Here $\gamma^- = N[\omega^p]$ is already algebraic (complete intersections of hyperplanes).

2. **Calibration–coercivity (unchanged).** For effective γ^+, the coercivity inequality forces energy-minimizing sequences toward the calibrated cone. This is the content of the earlier version you reviewed.

3. **SYR construction (now complete).** Build a sequence of integral ψ-calibrated cycles T_k with:

    - $\text{Mass}(T_k) \to m \int_X \beta \wedge \psi$,

- Tangent-plane Young measures converging to the prescribed distribution,
- $[T_k] = \mathrm{PD}(m[\gamma^+])$.

The new corner-exit vertex-template machinery provides the "$\mathcal{F}(\partial T^{\mathrm{raw}}) = o(m)$" estimate that was previously missing.

4. **Federer–Fleming compactness + Harvey–Lawson (standard).** The limit T is a ψ-calibrated integral current, hence a positive sum of complex analytic subvarieties. By Chow's theorem, these are algebraic.

5. **Conclusion.** $\gamma = [Z^+] - [Z^-]$ is algebraic.

5 What to Check During Review

For a careful review, I suggest focusing on:

1. `prop:corner-exit-template-net`: Does the perturbation argument correctly produce a finite net where every direction has corner-exit templates? The key is that the "bad" set (where some coefficient vanishes) is a finite union of proper algebraic subvarieties.

2. `lem:global-graph-contraction`: Is the contraction-mapping argument correctly set up? The domain must contain the cell Q, not just an infinitesimal neighborhood.

3. `lem:bergman-affine-approx-hormander`: The cutoff + Hörmander construction is standard, but verify that the exponential decay $O(e^{-cm})$ propagates correctly to C^1 estimates on the inner ball.

4. **Parameter regime consistency**: Multiple constraints must coexist:

   - $h \lesssim m^{-1/2}$ (Bergman control on cells),
   - $\delta \gtrsim \varepsilon h$ (disjointness of slivers),
   - $N_Q \gtrsim h^{-1}$ (many pieces for rounding),
   - The weighted flat-glue bound is $o(m)$.

   These are recorded in `rem:weighted-scaling`; verify they are mutually satisfiable.

5. **Slow variation + discrepancy rounding compatibility**: `lem:slow-variation-discrepancy` shows that 0–1 discrepancy rounding (Bárány–Grinberg style) preserves slow variation. Check that the "+2" error term doesn't break the $O(h)$-fraction bound.

6 Files in the Repository

- `hodge-SAVE-dec-12-handoff.tex` — the master manuscript (now unconditional)
- `mg-microstructure-gluing-target.txt` — dependency DAG for the microstructure step
- `strategy-and-progress.md` — lab notebook of the proof development
- `proof-completion-plan.md` — standalone completion plan
- `hodge-fix-dec-8-old.tex` — the earlier version you reviewed (calibration–coercivity)

Summary of Changes from the Previous Version

Previous Version	**Current Version**
Calibration–coercivity established	Unchanged
Quantitative approximation to calibrated cone	Unchanged
SYR realization flagged as "conditional on `rem:glue-gap`"	SYR realization is now unconditional
No explicit corner-exit geometry	Full corner-exit vertex-template construction
No Bergman-scale graph control	Global graph lemma + Hörmander construction
Headline theorems marked "Conditional"	Headline theorems are unconditional

Thank you for your continued review. Please contact me with any questions or concerns.